

Ex.No: 3 Date:19.03.2025	Accessing Text Corpora using NLTK in Python
---	--

AIM:

To access and work with available text corpora using NLTK libraries.

PROCEDURE:

1. Import required modules for text processing, visualization, and deep learning.
2. Load sample text data from the NLTK Gutenberg corpus.
3. Calculate lexical diversity by computing the ratio of unique words to total words.
4. Generate a word cloud to visualize the most frequent words in the text.
5. Tokenize the text into words and sentences using NLTK's tokenization functions.
6. Remove stopwords to filter out common but unimportant words.
7. Perform stemming and lemmatization to reduce words to their root forms.
8. Train a basic NLP model using TensorFlow to process text data.
9. Perform named entity recognition (NER) to identify proper names and entities.
10. Extract n-grams such as bigrams and trigrams from the text.

Program:**1. Import required modules.**

```
import nltk
from nltk.corpus import *
from nltk.tokenize import word_tokenize, sent_tokenize
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer, WordNetLemmatizer
from nltk.probability import FreqDist
nltk.download('all')
```

2. Load sample text data

```
nltk.download('gutenberg')
texts = gutenberg.fileids()
```

```
[ 'austen-emma.txt',
  'austen-persuasion.txt',
  'austen-sense.txt',
  'bible-kjv.txt',
  'blake-poems.txt',
  'bryant-stories.txt',
  'burgess-busterbrown.txt',
  'carroll-alice.txt',
  'chesterton-ball.txt',
  'chesterton-brown.txt',
  'chesterton-thursday.txt',
  'edgeworth-parents.txt',
  'melville-moby_dick.txt',
  'milton-paradise.txt',
  'shakespeare-caesar.txt',
  'shakespeare-hamlet.txt',
  'shakespeare-macbeth.txt',
  'whitman-leaves.txt']
```

```
gutenberg.words('shakespeare-hamlet.txt')[:50]
```

```
['[',
 'The',
 'Tragedie',
 'of',
 'Hamlet',
 'by',
 'William',
 'Shakespeare',
 '1599',
 ']',
 'Actus',
 'Primus',
 '.',
 'Scoena',
 'Prima',
 '.',
 'Enter',
 'Barnardo',
```

3. Calculate lexical diversity.

```
def lexical_diversity(text):
    words = gutenberg.words(text)
    unique_words = set(words)
    return len(unique_words) / len(words)

diversity_scores = {text: lexical_diversity(text) for text in texts}

sorted_diversity = sorted(diversity_scores.items(), key=lambda x: x[1], reverse=True)
for text, score in sorted_diversity:
    print(f'{text}: {score:.4f}')
```

```
blake-poems.txt: 0.2179
shakespeare-macbeth.txt: 0.1736
shakespeare-hamlet.txt: 0.1458
shakespeare-caesar.txt: 0.1378
milton-paradise.txt: 0.1110
chesterton-thursday.txt: 0.0983
chesterton-brown.txt: 0.0964
burgess-busterbrown.txt: 0.0930
whitman-leaves.txt: 0.0925
chesterton-ball.txt: 0.0922
carroll-alice.txt: 0.0884
bryant-stories.txt: 0.0795
melville-moby_dick.txt: 0.0741
austen-persuasion.txt: 0.0625
austen-sense.txt: 0.0483
edgeworth-parents.txt: 0.0455
austen-emma.txt: 0.0406
bible-kjv.txt: 0.0136
```

	Word	Frequency
0	ham	337
1	lord	211
2	haue	178
3	king	172
4	thou	107
5	shall	107
6	come	104
7	let	104
8	hamlet	100
9	good	98
10	hor	95
11	thy	90
12	enter	85
13	oh	81
14	like	80
15	would	73
16	well	71
17	know	71
18	tis	69
19	selfe	68

5. Text tokenization

```
tokenizer = Tokenizer()
tokenizer.fit_on_texts([hamlet_text])
total_words = len(tokenizer.word_index) + 1
```

6. LSTM Model for text generation

```
input_sequences = []
words = hamlet_text.split()
for i in range(3, len(words)):
    seq = words[i-3:i+1]
    input_sequences.append(tokenizer.texts_to_sequences([" ".join(seq))][0])
max_length = max([len(seq) for seq in input_sequences])
input_sequences = pad_sequences(input_sequences, maxlen=max_length, padding='pre')
X, y = input_sequences[:, :-1], input_sequences[:, -1]
y = tf.keras.utils.to_categorical(y, num_classes=total_words)
model = Sequential([
    Embedding(total_words, 100, input_length=max_length-1),
    LSTM(150, return_sequences=True),
    LSTM(150),
    Dense(150, activation='relu'),
    Dense(total_words, activation='softmax')
])
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(X, y, epochs=50, verbose=1)
def generate_shakespeare_text(seed_text, next_words=20):
    for _ in range(next_words):
```

```

token_list = tokenizer.texts_to_sequences([seed_text])[0]
token_list = pad_sequences([token_list], maxlen=max_length-1, padding='pre')
predicted = np.argmax(model.predict(token_list), axis=-1)
output_word = ""
for word, index in tokenizer.word_index.items():
    if index == predicted:
        output_word = word
        break
seed_text += " " + output_word
return seed_text
print(generate_shakespeare_text("to be or not to be"))

1/1 ————— 0s 80ms/step
1/1 ————— 0s 75ms/step
1/1 ————— 0s 70ms/step
1/1 ————— 0s 67ms/step
1/1 ————— 0s 69ms/step
1/1 ————— 0s 69ms/step
to be or not to be neere i businesse and disappointed pittious so puh how by our knotty and fallies but an masters countrey who whose

```

#MOVIE REVIEWS corpus:

1. Load and Preprocess Data

```

import nltk
nltk.download('movie_reviews')
from nltk.corpus import movie_reviews
len(movie_reviews.words())
movie_reviews.categories()
from collections import Counter
all_words=Counter(movie_reviews.words())

```

2.Feature Extraction

```

feature = {}
review = movie_reviews.words('neg/cv954_19932.txt')
for x in range(len(feature_vector)):
    feature[feature_vector[x]] = feature_vector[x] in review
[x for x in feature_vector if feature[x] == True]

```

3.Train Machine Learning Models

```

import nltk
import pandas as pd

```

```

import numpy as np
from nltk.corpus import movie_reviews
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score
import re
import string

nltk.download('movie_reviews')
nltk.download('punkt')

def load_movie_reviews():
    documents = []
    labels = []

    for category in movie_reviews.categories():
        for fileid in movie_reviews.fileids(category):
            documents.append(' '.join(movie_reviews.words(fileid)))
            labels.append(category)

    return documents, labels

def preprocess_text(text):
    text = text.lower()
    text = re.sub(f'[{string.punctuation}]', ' ', text)
    text = re.sub(r'\s+', ' ', text).strip()
    return text

documents, labels = load_movie_reviews()
processed_docs = [preprocess_text(doc) for doc in documents]
binary_labels = [1 if label == 'pos' else 0 for label in labels]
X_train, X_test, y_train, y_test = train_test_split(
    processed_docs, binary_labels, test_size=0.2, random_state=42
)
tfidf_vectorizer = TfidfVectorizer(max_features=5000, ngram_range=(1, 2))
X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)
X_test_tfidf = tfidf_vectorizer.transform(X_test)
sentiment_model = LogisticRegression(C=1.0, random_state=42)
sentiment_model.fit(X_train_tfidf, y_train)
y_pred = sentiment_model.predict(X_test_tfidf)
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred, target_names=['Negative', 'Positive']))

def predict_sentiment(new_reviews):
    processed_reviews = [preprocess_text(review) for review in new_reviews]
    X_new_tfidf = tfidf_vectorizer.transform(processed_reviews)
    predictions = sentiment_model.predict(X_new_tfidf)
    probabilities = sentiment_model.predict_proba(X_new_tfidf)

    results = []
    for i, pred in enumerate(predictions):

```

```

        sentiment = 'Positive' if pred > 0.5 else 'Negative'
        confidence = probabilities[i][pred]
        results.append({
            'review': new_reviews[i],
            'sentiment': sentiment,
            'confidence': confidence
        })

    return results

new_reviews = [
    "it is very good",
    "nice",
    "booring bad"
]
results = predict_sentiment(new_reviews)

results = predict_sentiment(new_reviews)
print("\nPredictions for New Reviews:")
for i, result in enumerate(results):
    print(f"\nReview {i+1}: {result['review']}")
    print(f"Predicted Sentiment: {result['sentiment']}")
    print(f"Confidence: {result['confidence']}")

```

Model Accuracy: 0.8400

Classification Report:

	precision	recall	f1-score	support
Negative	0.85	0.82	0.84	200
Positive	0.83	0.86	0.84	200
accuracy			0.84	400
macro avg	0.84	0.84	0.84	400
weighted avg	0.84	0.84	0.84	400

Predictions for New Reviews:

Review 1: it is very good
Predicted Sentiment: Positive
Confidence: 0.6004715807048746

Review 2: nice
Predicted Sentiment: Positive
Confidence: 0.5411814943329837

Review 3: booring bad
Predicted Sentiment: Negative
Confidence: 0.9528418976064467

Result:

Thus the above program accessing nltk corpora (Gutenberg and Movie Reviews) was executed successfully.